

# CSI3131 – Operating Systems

## Lab1

### Objectives:

Experiment with spawning processes using fork. Examining the process status as it progresses.

### Lab Setup

You will need a Unix/Linux OS shell with gcc/g++ c-compiler.

- If you're using MAC you're probably already setup for this Lab, as MAC OS is a UNIX based OS.
- If you're using Windows OS, you have many options:
  1. Windows 10 provides aSubsystem for linux, here is a link from Microsoft on how to set it up for windows 10: <https://docs.microsoft.com/en-us/windows/wsl/install-win10>
  2. Install Cygwin. Go to <http://preshing.com/20141108/how-to-install-the-latest-gcc-on-windows/> and follow the instructions to install Cygwin
  3. Use Linux server provided in the lab

### Part1

Systemverilog is hardware description language that is extensively used in ASIC/FPGA design and verification. The language supports the following three forms of the fork constructs for threads:

<b>fork</b> task1; task2; task3; <b>join</b>	<b>fork</b> task1; task2; task3; <b>join_none</b>	<b>fork</b> task1; task2; task3; <b>join_any</b>
--	---	--

In the above three fork scenarios, the process spawns 3 different threads of processing performing tasks in addition to the original thread that is also performing its own task.

Depending on the join construct used the behavior of the threads may differ as follows:

- The join construct makes the parent wait for all child threads to finish before resuming the parent process.
- The join\_none construct makes the parent not to wait for any child threads to finish before resuming the parent process, i.e. it runs concurrently with the child threads.

- The `join_any` construct makes the parent wait for at least one child process to finish before resuming the parent process.

Your task is to mimic the above behavior for processes rather than threads. Assuming you have 3 applications in the path named `task1`, `task2`, and `task3`.

The solution of the first mode is provided together with a task application. The attached `Lab1.zip` file includes a directory `part1` that contains 2 files:

1. `task.c` is a simple console application that sleeps for `n` seconds, marking the start/end of the process. The time period (`n`) is passed to the task as a command line argument
2. `fork_join.c` is the solution of the fork-join scenario as explained above. The `fork_join` application takes up to 4 arguments as command line arguments, namely `task1`, `task2`, `task3`, and parent task time periods, respectively. Checkout how the task is launched inside the child process.  
<https://linux.die.net/man/3/execvp>
3. Copy the `fork_join.c` file to `fork_join_none.c` and make the necessary modifications to make it behave as a fork-join-none case.
4. Copy the `fork_join.c` file to `fork_join_any.c` and make the necessary modifications to make it behave as a fork-join-any case.

NOTE: Keep the logging messages intact as it will be used to automatically mark your assignment. Failure to do so will result in automatic deductions!

Also attached is the expected results for few of the test cases that will be used in marking, other cases will also be used to verify the correctness of your code.

5. Your submission should include only the 4 c-files named above. Don't include executables not log files.

## Part2

Linux presents information on processes via the file system. Statistics on running processes can be examined by opening regular files. Note that these files do not exist on the hard drive, but in main memory. In fact they are not files at all, but simply system data presented as files.

The directory `/proc` contains a set of directories and files that provides access to general OS information about each process running on the system. Each process has a unique identifier, called the PID (Process Identifier). In `/proc`, the PID is presented as a directory. The contents of the PID directory provides information on the corresponding process. You can obtain documentation on the `proc` directory using "`man proc`" command.

In this part of the lab, you shall examine how a process executes and changes. Some utilities (shell programs) and C programs have been provided to examine the changes in processes (by reading the contents of the `/proc` directory) as they execute different programs.

You will gain insight into the state of a process and the difference between running in the two modes of execution (kernel and user).

Enter command: `ls /proc`. The content of the `/proc` directory will be listed. You will see many number entries (representing directories containing information about processes, for example directory 23406 contains information about process with PID 23406), as well as some more meaningful entries (for example 'version' and 'cpuinfo').

Enter command `cat /proc/version | more`. This should display the content of the file `/proc/version` that contains information about the OS version.

Enter command `cat /proc/cpuinfo | more`. You might want to explore the content of other files as well.

Enter command `ps`. This will list the processes you have launched. At least two processes will be listed: 'ps' (the one you just launched and which is producing this and a shell process (probably 'bash', but could be different, depending on your settings)). Write down the PID of the shell process.

Enter command `cat /proc/NNN/stat | more`, where NNN is the PID of your shell process from the previous step. You will see a bunch of numbers containing lots of information about this process. You can learn the meaning of the numbers by entering `man proc | more` command, which would show you the manual for `proc`. You can see most of this information in human readable form in the file `/proc/NNN/status`. Try `cat /proc/NNN/status`

1. Compile and run `calclloop.c` (`g++ calclloop.c -o calclloop.exe`)
2. Compile and run `cploop.c`
3. Compile `procmon.c` and preview its code – no need to modify this file
4. Run `calclloop.exe` in the background and run `procmon.exe` passing it `calclloop's` PID – checkout the process progress, it should look somewhat like the following:

```
$ ./calclloop &  
[1] 6740  
  
$ ./procmon 6740
```

```
Monitoring /proc/6740/stat:  
  
Time      State      SysTm     UstTm  
0         Sleeping(memory)  0         4062  
1         Running      0         4906  
2         Sleeping(memory)  0         5078  
3         Sleeping(memory)  0         5078  
4         Sleeping(memory)  0         5078  
5         Running      0         5937
```

```

6      Sleeping(memory)      0      6062
7      Sleeping(memory)      0      6062
8      Sleeping(memory)      0      6062
9      Running                0      6953
10     Sleeping(memory)      0      7015
11     Sleeping(memory)      0      7015
12     Sleeping(memory)      0      7015
13     Running                0      7984
14     Sleeping(memory)      0      8062
15     Sleeping(memory)      0      8062
16     Sleeping(memory)      0      8062
17     Running                0      9000
18     Sleeping(memory)      0      9031
19     Sleeping(memory)      0      9031
20     Sleeping(memory)      0      9031
procmon: Cannot open /proc/6740/stat, the monitored process is not running any more.
[1]+  Done                  ./calclloop

```

Notice how the user-time is incremented every time the process is running and remains constant every time the process is sleeping (idle)

5. Run cploop.exe in the background and run procmon.exe passing it cploop's PID – checkout the process progress, it should look somewhat like the following:

```

$ ./cploop.exe &
[1] 6884

$ ./procmon 6884

      Monitoring /proc/6884/stat:

Time      State      SysTm      UstTm
0      Running      1859      265
1      Sleeping(memory) 2093      343
2      Sleeping(memory) 2093      343
3      Sleeping(memory) 2093      343
4      Running      2656      484
5      Sleeping(memory) 2687      484
6      Sleeping(memory) 2687      484
7      Sleeping(memory) 2687      484
8      Sleeping(memory) 3312      625
9      Sleeping(memory) 3312      625
10     Sleeping(memory) 3312      625
11     Running      3515      656
12     Sleeping(memory) 3921      734
13     Sleeping(memory) 3921      734
14     Sleeping(memory) 3921      734
15     Running      4390      812
16     Sleeping(memory) 4546      843
17     Sleeping(memory) 4546      843
18     Sleeping(memory) 4546      843
19     Sleeping(memory) 5203      921
20     Sleeping(memory) 5203      921
21     Sleeping(memory) 5203      921
22     Running      5312      937
23     Sleeping(memory) 5890      1156
24     Sleeping(memory) 5890      1156
25     Sleeping(memory) 5890      1156
26     Running      6078      1203
27     Sleeping(memory) 6515      1281
28     Sleeping(memory) 6515      1281
29     Sleeping(memory) 6515      1281
30     Running      6968      1328
procmon: Cannot open /proc/6884/stat, the monitored process is not running any more.
[1]+  Done                  ./cploop.exe

```

Notice how both system and user times are incremented every time the process is running. Reflect on why this is the case.

6. Your task is to complete the code of mon.c, which is programmatically performing steps 1-5. Follow the instructions commented in the code:

```
/* Here comes your code.*/  
/* It should do the following:  
    1. fork/launch the program 'calclloop' and get its pid  
    2. fork/launch 'procmon pid' where pid is the pid of the process  
    launched in step 1  
    3. wait till calclloop process ends  
    4. wait till procmon process ends  
    5. fork/launch the program 'cploop' and get its pid  
    6. fork/launch 'procmon pid' where pid is the pid of the process  
    launched in step 5  
    7. wait till cploop process ends  
    8. wait till procmon process ends  
*/
```

NOTE: if you cannot launch procmon process twice you may split the job into two files: 'mon\_calclloop.c' & 'mon\_cploop.c'

7. Include only the source code for the files in part2 directory  
8. Zip part1 & part2 directory in one zip file lab1.zip and submit on bright space